

Final Project

Convolutional Networks for Image classification : Rock, Paper, Scissors

Maddie Lebiezinski and Ethan Kong

```
In [1]: 1 from tensorflow.keras.preprocessing.image import ImageDataGenerator
2 from tensorflow.keras.preprocessing import image
3 from tensorflow.keras.optimizers import RMSprop
4 import tensorflow as tf
5 import matplotlib.pyplot as plt
6 import os
7 import cv2
8 import numpy as np
```


2023-04-28 16:29:40.273966: I tensorflow/core/platform/cpu_feature_guard.cc:193] This TensorFlow binary is optimized with oneAPI Deep Neural Network Library (oneDNN) to use the following CPU instructions in performance-critical operations: SSE4.1 SSE4.2

To enable them in other operations, rebuild TensorFlow with the appropriate compiler flags.

```
In [2]: 1 # for train images
2 train_path = "Rock-Paper-Scissors/train/"
3 # for test images
4 test_path = "Rock-Paper-Scissors/test/"
5 # validation
6 validation_path = "Rock-Paper-Scissors/validation/"
```

Load and visualize data

```
In [3]: 1 # example of a loaded image
        2 img = image.load_img(train_path + "/paper/paper01-000.png")
        3 plt.imshow(img)
        4 plt.axis("on")
        5 plt.show()
```



```
In [4]: 1 # shape of this image
        2 cv2.imread(train_path + "/paper/paper01-000.png").shape
        3 # output : 300 pixels height, 300 pixels width, rgb 3 colors
```

Out[4]: (300, 300, 3)

```
In [5]: 1 # generating training
        2 train = ImageDataGenerator(rescale= 1/255) # because rgb values go up
        3 validation = ImageDataGenerator(rescale= 1/255)
```

In []: 1

```
In [6]: 1 train_dataset = train.flow_from_directory(train_path,
        2                                           target_size= (300, 300), # in
        3                                           batch_size = 32, #
        4                                           class_mode = 'categorical') #
        5
        6 validation_dataset = train.flow_from_directory(validation_path,
        7                                           target_size= (300, 300), # in
        8                                           batch_size = 32, #
        9                                           class_mode = 'categorical') #
```

Found 2520 images belonging to 3 classes.
Found 33 images belonging to 3 classes.

```
In [7]: 1 train_dataset.class_indices
```

Out[7]: {'paper': 0, 'rock': 1, 'scissors': 2}

```
In [8]: 1 train_dataset.classes
```

Out[8]: array([0, 0, 0, ..., 2, 2, 2], dtype=int32)

Construct deep learning model

```

In [9]: 1 model = tf.keras.models.Sequential([tf.keras.layers.Conv2D(16,(3, 3),
2                                           tf.keras.layers.MaxPool2D(2,2),
3
4                                           tf.keras.layers.Conv2D(32,(3, 3),
5                                           tf.keras.layers.MaxPool2D(2,2),
6
7                                           tf.keras.layers.Conv2D(64,(3, 3),
8                                           tf.keras.layers.MaxPool2D(2,2),
9
10                                          tf.keras.layers.Flatten(),
11
12                                          tf.keras.layers.Dense(2892, activation='relu',
13                                                                  # number of
14                                          tf.keras.layers.Dense(3, activation='softmax',
15                                                                  # 3 categories
16 ])
17
18 # to view all the layers of the network using the Keras Model.summary
19 model.summary()

```

2023-04-28 16:29:45.475462: I tensorflow/core/platform/cpu_feature_guard.cc:193] This TensorFlow binary is optimized with oneAPI Deep Neural Network Library (oneDNN) to use the following CPU instructions in performance-critical operations: SSE4.1 SSE4.2
To enable them in other operations, rebuild TensorFlow with the appropriate compiler flags.

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 298, 298, 16)	448
max_pooling2d (MaxPooling2D)	(None, 149, 149, 16)	0
conv2d_1 (Conv2D)	(None, 147, 147, 32)	4640
max_pooling2d_1 (MaxPooling2D)	(None, 73, 73, 32)	0
conv2d_2 (Conv2D)	(None, 71, 71, 64)	18496
max_pooling2d_2 (MaxPooling2D)	(None, 35, 35, 64)	0
flatten (Flatten)	(None, 78400)	0
dense (Dense)	(None, 2892)	226735692
dense_1 (Dense)	(None, 3)	8679

```

=====
Total params: 226,767,955
Trainable params: 226,767,955
Non-trainable params: 0
=====

```

```
In [10]: 1 # compile the model
          2 model.compile(loss= 'categorical_crossentropy',
          3                   optimizer = RMSprop(lr=0.001),
          4                   metrics = ['accuracy'])
```

WARNING:absl:`lr` is deprecated, please use `learning_rate` instead, or use the legacy optimizer, e.g.,`tf.keras.optimizers.legacy.RMSprop`.

Training

```
In [11]: 1 model_fit= model.fit(train_dataset,
2                               steps_per_epoch = 20,
3                               epochs= 20,
4                               validation_data = validation_dataset)
```

```
Epoch 1/20
20/20 [=====] - 75s 4s/step - loss: 2.1607 - acc
uracy: 0.5222 - val_loss: 0.8475 - val_accuracy: 0.7576
Epoch 2/20
20/20 [=====] - 73s 4s/step - loss: 0.9757 - acc
uracy: 0.7281 - val_loss: 0.7622 - val_accuracy: 0.6667
Epoch 3/20
20/20 [=====] - 71s 4s/step - loss: 0.2048 - acc
uracy: 0.9399 - val_loss: 0.9675 - val_accuracy: 0.7273
Epoch 4/20
20/20 [=====] - 71s 4s/step - loss: 0.0252 - acc
uracy: 0.9953 - val_loss: 1.1462 - val_accuracy: 0.6667
Epoch 5/20
20/20 [=====] - 65s 3s/step - loss: 0.0172 - acc
uracy: 0.9968 - val_loss: 1.0417 - val_accuracy: 0.7879
Epoch 6/20
20/20 [=====] - 75s 4s/step - loss: 0.4563 - acc
uracy: 0.8906 - val_loss: 0.7344 - val_accuracy: 0.7576
Epoch 7/20
20/20 [=====] - 71s 4s/step - loss: 0.0166 - acc
uracy: 0.9984 - val_loss: 0.8632 - val_accuracy: 0.8182
Epoch 8/20
20/20 [=====] - 68s 3s/step - loss: 0.0163 - acc
uracy: 0.9953 - val_loss: 1.1737 - val_accuracy: 0.7576
Epoch 9/20
20/20 [=====] - 77s 4s/step - loss: 8.2154e-04 -
accuracy: 1.0000 - val_loss: 1.3731 - val_accuracy: 0.7576
Epoch 10/20
20/20 [=====] - 70s 3s/step - loss: 3.9407e-04 -
accuracy: 1.0000 - val_loss: 1.5656 - val_accuracy: 0.7273
Epoch 11/20
20/20 [=====] - 62s 3s/step - loss: 1.7029 - acc
uracy: 0.9019 - val_loss: 1.0658 - val_accuracy: 0.7576
Epoch 12/20
20/20 [=====] - 68s 3s/step - loss: 0.0054 - acc
uracy: 1.0000 - val_loss: 1.6269 - val_accuracy: 0.6970
Epoch 13/20
20/20 [=====] - 101s 5s/step - loss: 0.0013 - ac
curacy: 1.0000 - val_loss: 1.4174 - val_accuracy: 0.7273
Epoch 14/20
20/20 [=====] - 85s 4s/step - loss: 8.8733e-04 -
accuracy: 1.0000 - val_loss: 1.6309 - val_accuracy: 0.7273
Epoch 15/20
20/20 [=====] - 83s 4s/step - loss: 1.9818e-04 -
accuracy: 1.0000 - val_loss: 1.6107 - val_accuracy: 0.7273
Epoch 16/20
20/20 [=====] - 84s 4s/step - loss: 0.9632 - acc
uracy: 0.9225 - val_loss: 0.6992 - val_accuracy: 0.8182
Epoch 17/20
20/20 [=====] - 88s 4s/step - loss: 0.0051 - acc
uracy: 1.0000 - val_loss: 0.7134 - val_accuracy: 0.8182
Epoch 18/20
```


```
20/20 [=====] - 88s 4s/step - loss: 0.0023 - accuracy: 1.0000 - val_loss: 0.7450 - val_accuracy: 0.8182
Epoch 19/20
20/20 [=====] - 81s 4s/step - loss: 0.0013 - accuracy: 1.0000 - val_loss: 0.9910 - val_accuracy: 0.8182
Epoch 20/20
20/20 [=====] - 70s 3s/step - loss: 4.3057e-04 - accuracy: 1.0000 - val_loss: 1.1472 - val_accuracy: 0.8182
```

```
In [17]: 1 model_fit.history.keys()
```


```
Out[17]: dict_keys(['loss', 'accuracy', 'val_loss', 'val_accuracy'])
```

Performance plot

```
In [25]: 1 fig = plt.figure()
2 loss = model_fit.history['loss']
3 val_loss = model_fit.history['val_loss']
4 plt.plot(loss, color='red', label='loss')
5 plt.plot(val_loss, color='blue', label='val_loss')
6 fig.suptitle('Loss', fontsize=20)
7 plt.legend(loc="upper left")
8 plt.show()
```



```
In [22]: 1 fig = plt.figure()
2 acc = model_fit.history['accuracy']
3 val_acc = model_fit.history['val_accuracy']
4 plt.plot(acc, color='red', label='accuracy')
5 plt.plot(val_acc, color='blue', label='val_accuracy')
6 fig.suptitle('Accuracy', fontsize=20)
7 plt.legend(loc="upper left")
8 plt.show()
```



```
In [23]: 1 validation_dataset.class_indices
```


```
Out[23]: {'paper': 0, 'rock': 1, 'scissors': 2}
```

Test


```

In [24]: 1 # testset
          2 dir_path = 'Rock-Paper-Scissors/test_all'
          3
          4 for i in os.listdir(dir_path):
          5     img = image.load_img(dir_path+'/' + i, target_size= (300, 300))
          6     plt.imshow(img)
          7     plt.show()
          8
          9     X = image.img_to_array(img)
         10     X = np.expand_dims(X, axis =0)
         11     images = np.vstack([X])
         12     val = model.predict(images)
         13     print(val)
         14     if val[0][1] == 1:
         15         print("rock")
         16     elif val[0][0] == 1:
         17         print ("paper")
         18     elif val[0][2] == 1:
         19         print ("scissors")

```



```

1/1 [=====] - 3s 3s/step
[[0. 0. 1.]]
scissors

```

Link to dataset : <https://www.kaggle.com/datasets/sanikamal/rock-paper-scissors-dataset>
[\(https://www.kaggle.com/datasets/sanikamal/rock-paper-scissors-dataset\)](https://www.kaggle.com/datasets/sanikamal/rock-paper-scissors-dataset)